



BC2406 Business Analytics I: Predictive Techniques

Seminars 8

Classification and Decision Tree II

Instructor: Prof. Lee Gun-woong

Nanyang Business School

Review & Supplementary Slides


Decision Trees

Tid	Attrib1	Attrib2	Attrib3	Class
1	Yes	Large	125K	No
2	No	Medium	100K	No
3	No	Small	70K	No
4	Yes	Medium	120K	No
5	No	Large	95K	Yes
6	No	Medium	60K	No
7	Yes	Large	220K	No
8	No	Small	85K	Yes
9	No	Medium	75K	No
10	No	Small	90K	Yes

Training Set

Tid	Attrib1	Attrib2	Attrib3	Class
11	No	Small	55K	?
12	Yes	Medium	80K	?
13	Yes	Large	110K	?
14	No	Small	95K	?
15	No	Large	67K	?

Test Set



Tree Construction

- Build a Decision Tree
 - Choose the *best* attribute(s) to split the remaining instances and make that attribute a decision node
 - Select the attribute that produces the “**purest**” nodes
 - Repeat this process recursively for each child node
 - Stop when:
 - All (almost all) the instances have the same class attribute value
 - There are no more instances / attributes
- Determine how to split the instances
 - How to determine the best split?
 - Nodes with **homogeneous** class distribution are preferred

C0: 5
C1: 5

Non-homogeneous,
High degree of impurity

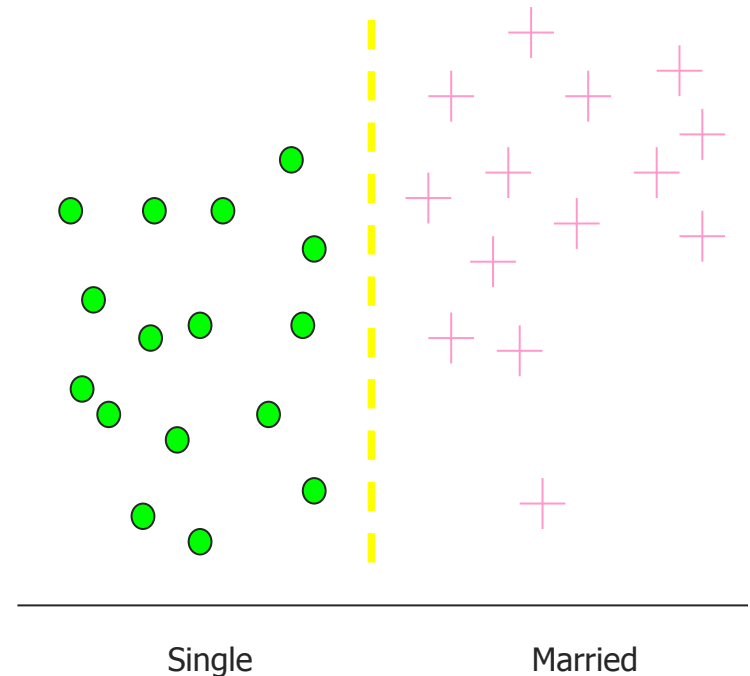

C0: 9
C1: 1

Homogeneous,
Low degree of impurity

Impurity

Which attributes is more informative (purer)?

Split over whether income exceeds 50K **Split over whether an applicant is married**



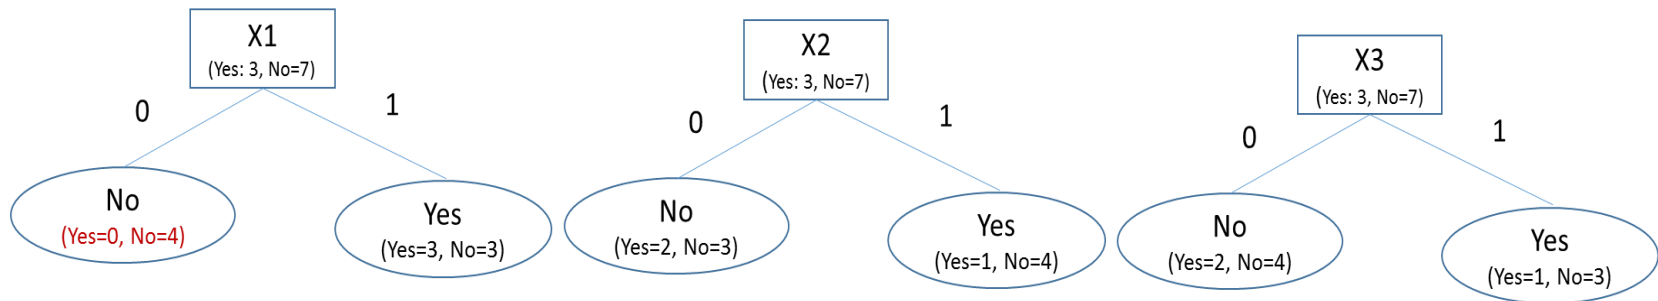
Example: Tree Construction

- Consider the training data given below. In the dataset, X1, X2, X3 are the attributes and Y is the class attribute.
- Draw the (full) decision tree for this dataset using the Gini Index

X1	X2	X3	Y
1	0	0	Yes
1	0	0	No
1	0	1	Yes
1	1	0	No
1	1	0	Yes
1	1	1	No
0	0	0	No
0	1	0	No
0	0	1	No
0	1	1	No

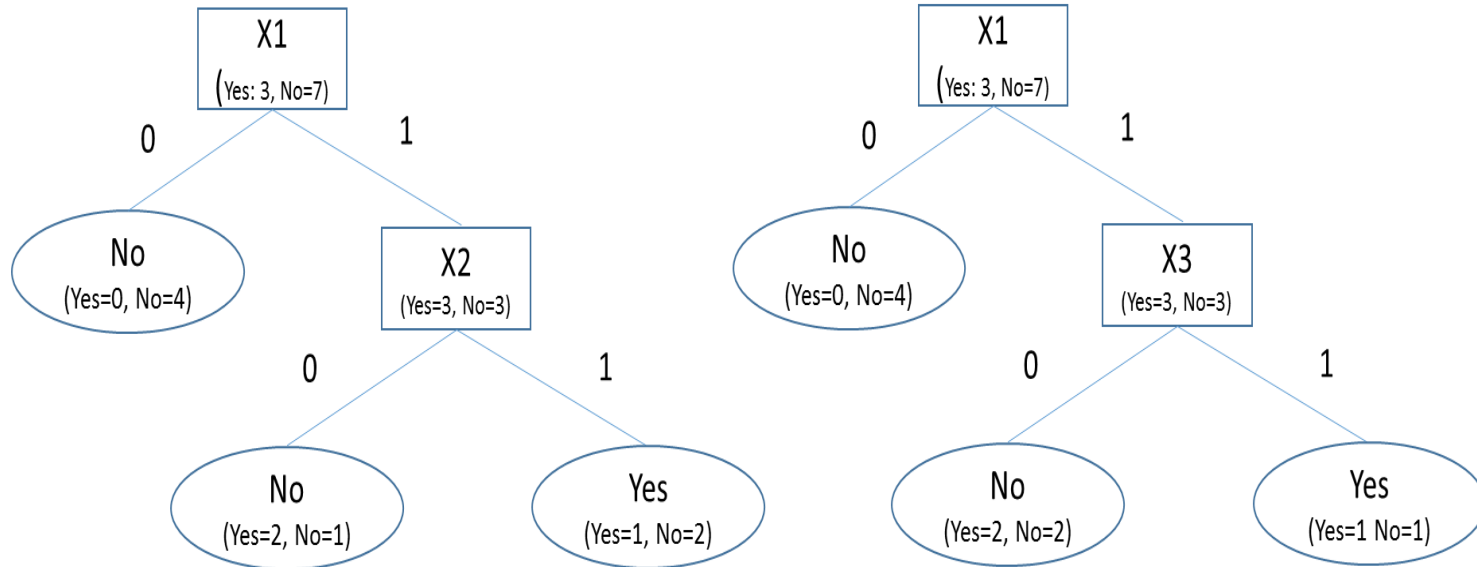
Example: Tree Construction (Cont.)

- At the first node:
 - Before Splitting
 - $\text{Gini}(Y) = 1 - P(\text{Yes})^2 - P(\text{No})^2 = 1 - (0.3)^2 - (0.7)^2 = 0.42$
 - After Splitting
 - $\text{Gini}(X1) = 6/10 * \text{Gini}(1) + 4/10 * \text{Gini}(0) = 0.6 * 0.5 + 0.4 * 0 = 0.30$
 - $\text{Gini}(X2) = 5/10 * \text{Gini}(1) + 5/10 * \text{Gini}(0) = 0.5 * 0.32 + 0.5 * 0.48 = 0.40$
 - $\text{Gini}(X3) = 4/10 * \text{Gini}(1) + 6/10 * \text{Gini}(0) = 0.4 * 0.375 + 0.6 * 0.44 = 0.42$



Example: Tree Construction (Cont.)

- At the second node:
 - Before Splitting
 - $\text{Gini}(Y|X1=1) = 1 - P(\text{Yes})^2 - P(\text{No})^2 = 1 - (0.5)^2 - (0.5)^2 = 0.5$
 - After Splitting
 - $\text{Gini}(X2|X1=1) = 3/6 * \text{Gini}(1|X1=1) + 3/6 * \text{Gini}(0|X1=1) = 0.5 * 0.44 + 0.5 * 0.44 = 0.44$
 - $\text{Gini}(X3|X1=1) = 2/6 * \text{Gini}(1|X1=1) + 4/6 * \text{Gini}(0|X1=1) = 0.5 * 0.50 + 0.5 * 0.50 = 0.50$



Example: Tree Construction (Cont.)

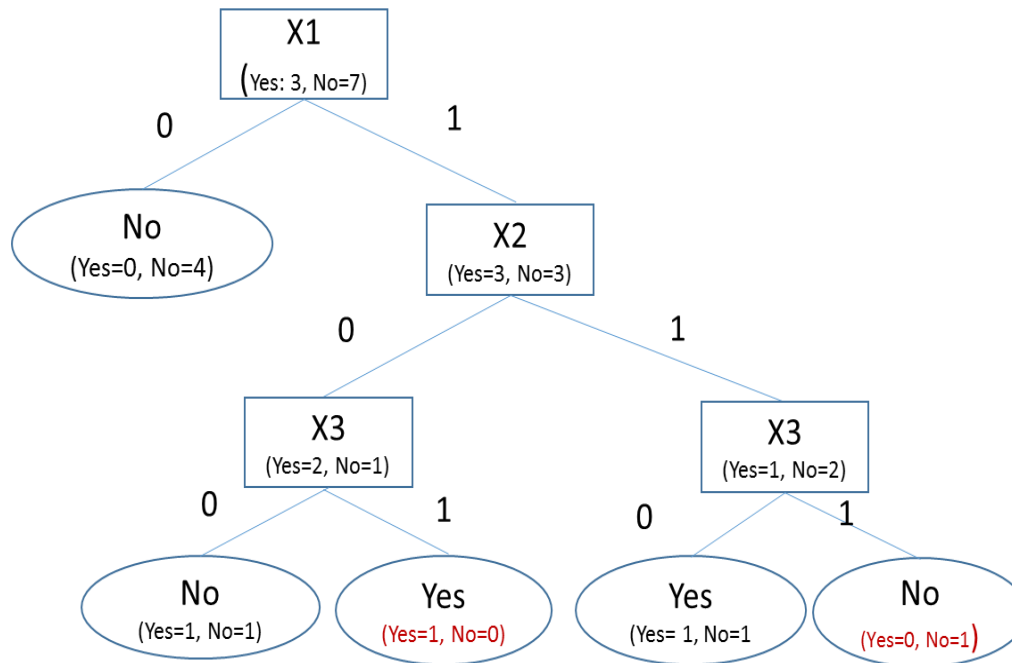
- At the Third node:

- Before Splitting

- $\text{Gini}(Y|X1=1, X2=0) = 1 - P(\text{Yes})^2 - P(\text{No})^2 = 1 - (2/3)^2 - (1/3)^2 = 0.44$
 - $\text{Gini}(Y|X1=1, X2=1) = 1 - P(\text{Yes})^2 - P(\text{No})^2 = 1 - (1/3)^2 - (2/3)^2 = 0.44$

- After Splitting

- $\text{Gini}(X3|X1=1, X2=0) = 2/3 * \text{Gini}(1|X1=1, X2=0) + 1/3 * \text{Gini}(0|X1=1, X2=0) = 1/3 * 0.0 + 2/3 * 0.5 = 0.33$
 - $\text{Gini}(X3|X1=1, X2=1) = 1/3 * \text{Gini}(1|X1=1, X2=1) + 2/3 * \text{Gini}(0|X1=1, X2=1) = 1/3 * 0.0 + 2/3 * 0.5 = 0.33$



Binary Attributes: GINI Index

- Splits into two partitions
- Finds larger and purer partitions

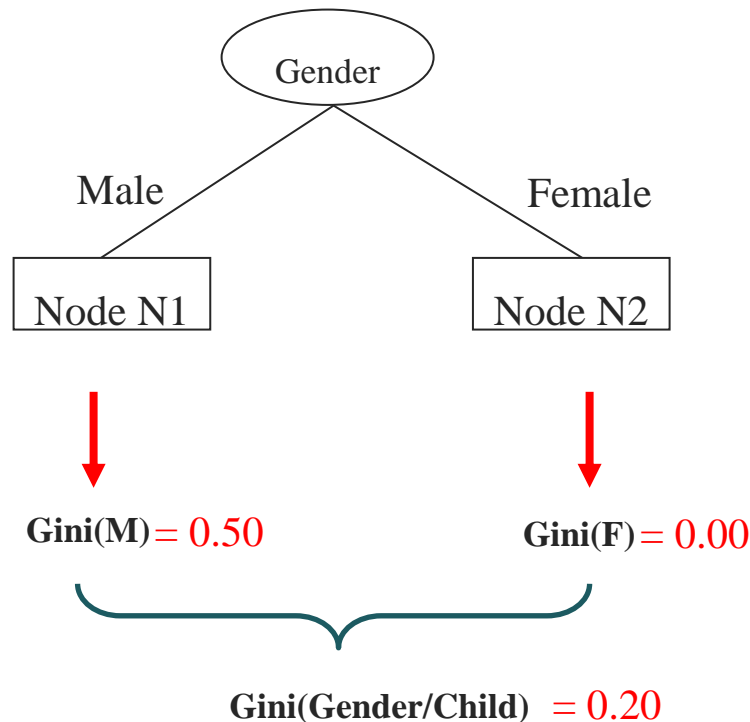
Before Splitting:

Yes
No



$\text{Gini}(\text{Before/Parent}) = 0.32$

After Splitting:



Categorical Attributes: Gini Index

- For each distinct value, gather counts for each class in the dataset
- Use the count matrix to make decisions

Multi-way split

	Marital Status		
	Single	Married	Divorced
Yes	1	2	1
No	4	1	1
Gini	0.393		

Two-way split
(find best partition of values)


	Marital Status	
	{Married, Divorced}	{Single}
Yes	3	1
No	2	4
Gini	0.400	

	Marital Status	
	{Married}	{Single, Divorced}
Yes	2	2
No	1	5
Gini	0.419	

Continuous Attributes: Gini Index

- Use Binary decisions based on one value
- Several Choices for the splitting value
 - Number of possible splitting values
= Number of distinct values
- Each splitting value (v) has a count matrix associated with it
 - Class counts in each of the partitions
 - $Income < v$ and $Income \geq v$
- Simple method to choose best v
 - For each v , scan the database to gather count matrix and compute its Gini index
 - Computationally Inefficient! Repetition of work.

Tid	Refund	Marital Status	Taxable Income	Cheat
1	Yes	Single	125K	No
2	No	Married	100K	No
3	No	Single	70K	No
4	Yes	Married	120K	No
5	No	Divorced	95K	Yes
6	No	Married	60K	No
7	Yes	Divorced	220K	No
8	No	Single	85K	Yes
9	No	Married	75K	No
10	No	Single	90K	Yes



Continuous Attributes: Gini Index...

- For efficient computation: for each attribute,
 - Sort the attribute on values
 - Linearly scan these values, each time updating the count matrix and computing gini index
 - Choose the split position that has the least gini index

		Cheat	No		No		No		Yes		Yes		Yes		No		No		No		No			
			Taxable Income																					
Sorted Values	→		60		70		75		85		90		95		100		120		125		220			
Split Positions	→		55		65		72		80		87		92		97		110		122		172		230	
			<=	>	<=	>	<=	>	<=	>	<=	>	<=	>	<=	>	<=	>	<=	>	<=	>	<=	>
		Yes	0	3	0	3	0	3	0	3	1	2	2	1	3	0	3	0	3	0	3	0	3	0
		No	0	7	1	6	2	5	3	4	3	4	3	4	3	4	4	3	5	2	6	1	7	0
		Gini	0.420		0.400		0.375		0.343		0.417		0.400		<u>0.300</u>		0.343		0.375		0.400		0.420	

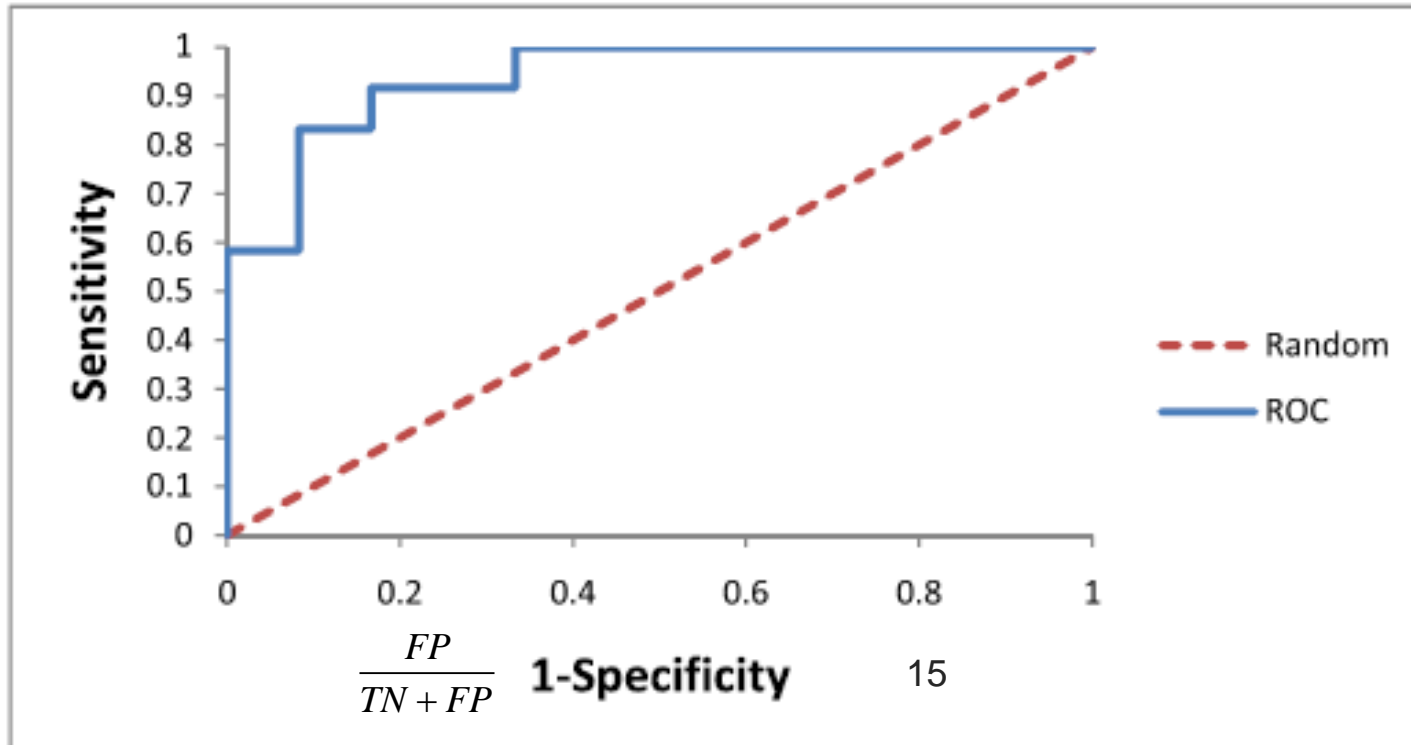
Model Evaluation

ACTUAL CLASS	PREDICTED CLASS		
		Class=1 (Important)	Class=0 (Less Important)
	Class=1 (Important)	a (20) (True Positive)	b (30) (False Negative)
	Class=0 (Less Important)	c (50) (False Positive)	d (100) (True Negative)

- Accuracy = $\frac{a+d}{a+b+c+d} = \frac{20+100}{20+30+50+100} = 0.6$
- Error Rate = 1- Accuracy = 0.4
- Sensitivity (True Positive Rate) = $\frac{a}{a+b} = \frac{20}{20+30} = 0.4$
 - Accuracy in classifying the important class correctly
- Specificity (True Negative Rate) = $\frac{d}{c+d} = \frac{100}{50+100} = 0.67$
 - Accuracy in classifying the less important class correctly


ROC Curve

- Receiver Operating Characteristic
- The closer to upper left corner the better



Overfitting and Pruning

Overfitting



- A good classification model must not only fit the training data well, it must also accurately classify instances it has never seen before.

– Low Training Error and Low Generalisation Error

Notes on Overfitting


- Overfitting results in decision trees that are more complex than necessary
 - Needs a shallower tree
- Training error no longer provides a good estimate of how well the tree will perform on previously unseen records
 - Improves generalization error

How to avoid Overfitting...

- Pruning (**Optimization**) : Identify and remove nodes in the decision tree that are not useful for classification
- **Post-pruning**
 - Grow a decision tree entirely that over-fits the training data
 - Evaluate the tree using the **test data**
 - Remove the nodes that have little effect on the classification errors
 - If classification error improves after trimming, replace sub-tree by a leaf node.
 - Stop when no more improvement
 - Use a set of different data from the training data to decide which is the —best pruned tree


Subtree Replacement

- Entire subtree is replaced by a single leaf node



Subtree Replacement

- Node 6 replaced the subtree
- Generalises tree a little more, but may increase accuracy





Cost Complexity Pruning



- Minimise the cost in misclassifications
 - $C_\alpha(T) = R(T) + \alpha |T|$
 - $C_\alpha(T)$: total cost in misclassifications
 - $R(T)$: the expected misclassification rate
 - the fraction of misclassified instances
 - $|T|$: the number of terminal nodes in the tree
 - α : a cost complexity parameter (CP)
 - $\alpha = 0$, the original fully grown tree
 - $\alpha = \infty$, the model with no splits at all (i.e., a single node)
 - Smaller CP means larger decision tree
 - As α increases, we incur a penalty that is proportional to the number of terminal nodes.
 - This will cause the minimum cost to occur for a tree that is a subtree of the original one (since a subtree will have a smaller number of terminal nodes).
 - Vary α and pick the value that gives the subtree that results in the smallest prediction error.

Cost Complexity Pruning



Full Tree



Start Pruning



Too Much!



Pruning with Re-sampling

- Resampling the training data allows us to know when we can make the best choice for the values of complexity parameter (α).
- Why do we need resampling approaches?
 - Comparing classifiers (rules) for the given dataset
 - Different classifiers will favor different domain of datasets
 - One needs to estimate how good the prediction will be.
- Resampling and Accuracy Measures
 - **Holdout** – randomly partition the given data into two independent sets and use one for training (e.g., 80%) and the other for testing (e.g., 20%)
 - **K-fold cross-validation** – randomly partition the given data into k-mutually exclusive subsets (folds). Training and testing is performed k times.


K-Fold Cross-Validation

- Misclassification errors are obtained through cross-validation to select the best model
- **2-Fold Cross-Validation**
 - Partition the data into two equal-size subsets (S1 and S2)
 - Run1: choose S1 for training and S2 for testing
 - Run2: choose S2 for training and S1 for testing
 - Average Error = (Error from Run1 + Error from Run2) / 2
- **K-fold Cross-Validation**
 - Segment the data into k equal-sized partitions
 - In each run, one of the partitions is chosen for testing, while others of them are used for training
 - Average Error = sum of the errors for all K runs / K
- **Repeated K-fold Cross-validation**
 - Creates multiple versions of the folds and aggregate the results


5-fold Cross-Validation




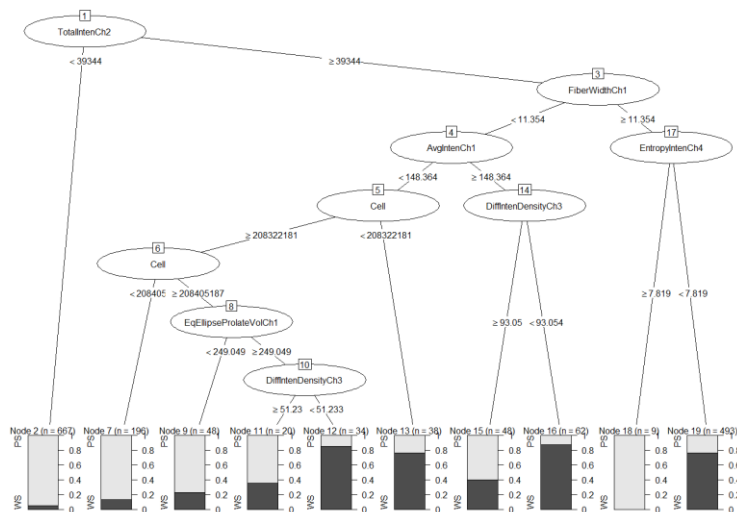
Cross-Validation: Example



- 10-fold cross validation with 3 repeats
 - Model performance over various levels of CP



- Before Pruning vs. After Pruning



Improving Model Performance

Accuracy Improvement

- **Bagging (Bootstrap aggregation)** – Number of trees are constructed on subsets of given data and majority voting is taken from these trees to classify a test sample.
- **Boosting** – attaching weights (importance) to the training samples and optimizing the weights during training and further using these weights to classify the test sample.

Boosting

- A method to “boost” weak learning algorithms (e.g., single trees) into strong learning algorithms.
- Boosting is an **ensemble-based method**.
 - It combines a number of weak learning algorithms to create a stronger learning algorithm than any of the algorithms alone.
 - The second tree corrects for the errors of the first tree, the third tree corrects for the errors of the first and second trees, and so forth.
 - Predictions are based on the entire ensemble of trees together that makes the prediction.
- **Learn in iterations**
 - Each iteration focuses on hard to learn parts of the attributes
 - i.e., instances that were misclassified by previous trees (weak learners)

Boosting Procedure

- Suppose the class attributes includes 20 instances
- 1st weak learner misclassifies 3 instances (D5, D10, D13)
- Update weights D_i
 - Weights of instances D5, D10 and D13 increase
 - Weights of other (correctly classified) instances decrease
- 2nd weak learner focuses more on the instances incorrectly classified by 1st weak learner and correctly classifies D10.
 - Weights of instances D5 and D13 increase
 - Weights of other (correctly classified) instances decrease
- 3rd weak learner focuses more on the instances incorrectly classified by 2nd weak learner...

Making Mistakes More Costlier than Others

- The cost of making a misclassification error may be higher for one class than the other(s)
- Looked at another way, the benefit of making a correct classification may be higher for one class than the other(s)
- Example: Identify risky bank loans (**default** or **not default**)
 - Giving a loan out to an applicant who is likely to default can be an expensive mistake.
 - The interest the bank would earn from a risky loan is far outweighed by the massive loss it would incur if the money is not paid back.

Making Mistakes More Costlier than Others

- Assign a penalty to different types of error in order to discourage a tree from making more costly mistakes
 - Reduce the number of false negatives (FN)
- The penalties are designated in a **cost matrix**, which specifies how much costlier each error is relative to any other prediction
- Cost Matrix Example
 - Suppose that a loan default costs the bank **four times** as much as a missed opportunity

	PREDICTED CLASS		
		Default	Not Default
	Default	0 (TP)	4 (FN)
ACTUAL CLASS	Not Default	1 (FP)	0 (TN)

Oversampling

- Asymmetric costs/benefits typically go hand in hand with presence of rare but important class
 - Responder to mailing
 - Someone who commits fraud
 - Debt defaulter
- Often we oversample rare cases to give model more information to work with

An Oversampling Procedure

- Typically use 50% “1” and 50% “0” for training
 - Method 1
 - Separate the responders (rare) from non-responders
 - Randomly select n responders, plus equal number of non-responders
 - Method 2
 - Replicating the existing class1’s several times to have n responders
 - Plus equal number of non-responders

Evaluation

- Evaluate the model on a testing set that has been selected without over- sampling (i.e., via simple random sampling).
- Evaluate the model on an oversampled testing set, and reweight the results to remove the effects of oversampling.